Attention Is All You Need

Ashish Vaswani, Noam Shazeer, Niki Parmar et al.

Synthetic fixture for unit tests (not the full paper).

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. Attention Is All You Need Ashish Vaswani, Noam Shazeer, Niki Parmar et al.

Attention Is All You Need — page 2

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 3

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 4

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 5

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 6

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 7

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 8

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 9

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 10

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 11

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.

Attention Is All You Need — page 12

The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings. The Transformer architecture relies on Attention mechanisms instead of recurrence. This synthetic page exists so CI can load a multi-page PDF without shipping the full arXiv article. Repeated filler about encoder layers, d_model, and positional encodings.